

1. Short Answer Questions

- **Q1:** Explain the primary differences between TensorFlow and PyTorch.

When would you choose one over the other?

- a) Execution Model
 - TensorFlow: Static computation graph (define-then-run)
 - PyTorch: Dynamic computation graph (define-by-run, more intuitive)
- b) Debugging
 - TensorFlow: Harder to debug due to static graph
 - PyTorch: Easier debugging with native Python tools (e.g., print, pdb)
- c) Deployment
 - TensorFlow: Better for production, mobile, and cross-platform deployment
 - PyTorch: Improving, but traditionally more research-focused
- d) Community & Adoption
 - TensorFlow: Backed by Google, widely used in industry
 - PyTorch: Backed by Meta, dominant in academia and research

When to Choose

- Choose TensorFlow:
 - Need production-ready tools
 - Targeting mobile/web deployment
 - Using Google Cloud ecosystem
- Choose PyTorch:
 - Rapid prototyping
 - Research and experimentation
 - Prefer Pythonic, flexible code
- **Q2:** Use Cases for Jupyter Notebooks in AI Development
 1. **Interactive Experimentation:** Ideal for testing models, visualizing data, and tweaking hyperparameters in real-time.
 2. **Documentation and Tutorials:** Combines code, output, and markdown to create readable, shareable learning materials or project walkthroughs.

- Q3: spaCy vs Basic Python String Operations in NLP

spaCy enhances NLP tasks by offering:

- **Tokenization, POS tagging, and Named Entity Recognition (NER)** out of the box, using trained statistical models.
- **Language-specific pipelines** that understand context, grammar, and semantics—far beyond simple `.split()` or regex.

2. Comparative Analysis

a) Target Application

i. Scikit- Learn

- Best for classical ML (e.g., regression, classification, clustering)
- Ideal for structured/tabular data.
- Not designed for deep learning

ii. Tensorflow:

- Built for deep learning (e.g., CNNs, RNNs, transformers).
- Suited for unstructured data (images, text, audio)
- Scales well for large datasets and production.

b) Ease of Use for Beginners

i. Scikit-Learn:

- Simple, consistent API (`fit`, `predict`, `transform`).
- Minimal setup, fast to prototype.
- Great for learning ML fundamentals.

ii. Tensorflow:

- Steeper learning curve (especially low-level API).
- Easier with Keras (high-level wrapper).
- Requires understanding of tensors and computational graphs.

c) Community Support.

iii. Scikit-Learn:

- Strong academic and data science community).
- Extensive documentation and tutorials.
- Stable and mature library.

iv. Tensorflow:

- Massive global community, backed by Google.
- Rich ecosystem (TensorBoard, TFLite, TF Hub)
- Frequent updates, large number of learning resources.

Part 2: Practical Implementation (50%)

Handwritten Model

```
# Model Evaluation
model.evaluate(X_test, y_test)
[65]

313/313 ━━━━━━━━ 1s 3ms/step - accuracy: 0.9785 - loss: 0.0734
[0.07338853925466537, 0.9785000085830688]
```

Iris Model

```
# Step 5: Evaluate the Model
[26]

# Evaluate using accuracy, precision, and recall
print("Accuracy:", accuracy_score(y_test, y_pred))
print("Precision (macro):", precision_score(y_test, y_pred, average='macro'))
print("Recall (macro):", recall_score(y_test, y_pred, average='macro'))
[27]

Accuracy: 1.0
Precision (macro): 1.0
Recall (macro): 1.0
```

NER results

```
139209 I just discovered Coconut Oil this year. Coco...
                                                Entities Sentiment
18828                                         []
363857                                         []
342609 [(Native Forest, ORG)] Neutral
62213                                         []
467133                                         []
518641 [(Delivery, ORG), (Connoisseur, ORG), (USDA, O... Positive
539725                                         [] Positive
```

Part 3: Ethics & Optimization

1. Ethical Considerations

a) Potential Biases

MNIST Model

- Class imbalance: Some digits (e.g., ‘1’, ‘7’) may appear more frequently, leading to skewed predictions.
- Handwriting style bias: Model may perform poorly on digits written in styles not well represented in the training set (e.g., regional or age-based variations).

Amazon Reviews Model

- Sentiment bias: Overrepresentation of positive or negative reviews can skew predictions.
- Demographic bias: Language patterns from certain regions or dialects may be misclassified.
- Product category bias: Sentiment may vary by category (e.g., electronics vs. books), but model may not account for this.

b) Mitigation Tools

TensorFlow Fairness Indicators

- Analyze model performance across slices (e.g., gender, age, product category).
- Identify disparities in accuracy, precision, recall across subgroups.
- Visualize fairness metrics to guide retraining or data augmentation.

spaCy’s Rule-Based Systems

- Customize token patterns to reduce bias in entity recognition or sentiment tagging.
- Apply consistent rules across dialects or product categories to normalize interpretation.
- Combine with statistical models to override biased predictions in edge cases

2. Troubleshoot Challenge

Sample Buggy Code

```
model = tf.keras.Sequential([
    tf.keras.layers.Dense(128, input_shape=(28, 28)),
    tf.keras.layers.Dense(10, activation='softmax')
])

model.compile(optimizer='adam', loss='binary_crossentropy', metrics=['accuracy'])
```

Issues:

- Input shape mismatch: input_shape=(28, 28) expects 2D input, but Dense layers need 1D.
- Incorrect loss function: binary_crossentropy is wrong for multi-class classification

Fixed Code:

```
model = tf.keras.Sequential([
    tf.keras.layers.Flatten(input_shape=(28, 28)), # Flatten 2D input
    tf.keras.layers.Dense(128, activation='relu'),
    tf.keras.layers.Dense(10, activation='softmax') # 10 classes
])
```

```
model.compile(optimizer='adam', loss='sparse_categorical_crossentropy', metrics=['accuracy'])
```

Explanation

Flatten() converts 28x28 input into 784-dim vector

'sparse_categorical_crossentropy' is correct when labels are integers (0–9).